

| | | |
|-------------------|---|--|
| EDUCATION | University of Copenhagen <i>Ph.D. in Natural Language Processing and Machine Learning</i> <ul style="list-style-type: none"> • Advisor: Prof. Anders Søgaard • Research area: Interpretability, reliability, and knowledge representations | Copenhagen, Denmark Jul 2021 - May 2025 |
| | University of Chile <i>Computer Science Engineer</i> <ul style="list-style-type: none"> • Cum. Grade: 6.5/7.0 (GPA 4) | Santiago, Chile March 2013 - July 2019 |
| EXPERIENCE | Google DeepMind <i>Student Researcher</i> <ul style="list-style-type: none"> • Led experiments on generating text attributed (citing) to source documents. • Trained T5-XXL models to generate text with planning and attribution. • Led the preparation of the data between 3 people constructing the pipeline. • Defined and developed the evaluation pipeline: text and attribution quality, and the relevance of the plan. | London, United Kingdom Oct 2022 – March 2023 |
| | Apple (Siri Understanding) <i>Research Intern</i> <ul style="list-style-type: none"> • Analyzed the performance of the latest multilingual parser model. • Proposed and experimented on improvements to the model by training modular language components. • Improved the performance of the model on low resource languages. | Cambridge, United Kingdom June 2022 – Sept 2022 |
| | Google (Youtube Video Classification) <i>Software Engineer</i> <ul style="list-style-type: none"> • Youtube classification platform for training, evaluating and governing classifiers. • Designed and implemented pipelines for automatic sampling, human labeling, training, and evaluation of classifiers. • In charge of building and maintaining 3 different binary classification models. • Developed evaluation analysis for the classifiers and collecting data pipelines. | Paris, France Oct 2019 – June 2021 |
| | How Do Multilingual Language Models Remember? <ul style="list-style-type: none"> • Mechanistic interpretability analyses of knowledge recall in multilingual LMs. • Addressing: Do the same mechanisms drive recall in different languages? and, when in the recalling process does language play a role? | Pre-print |
| RESEARCH PROJECTS | Defining Knowledge: Bridging Epistemology and LLMs <ul style="list-style-type: none"> • Addressed what it means to say that an LLM knows a propositional fact. • Surveyed epistemological definitions and formalized interpretations for LLMs. | EMNLP 2024 |
| | Learning to Plan and Generate Text with Citations <ul style="list-style-type: none"> • Internship at Google DeepMind (See above). | ACL 2024 |
| | MuLan: A Study of Fact Mutability in Language Models <ul style="list-style-type: none"> • Created a dataset and investigated if LLMs recognize that facts change over time. • Implemented minimum description length to probe the latent representations of the models for fact mutability. • Tested the accuracy of predictions and updates for different mutability levels. | NAACL 2024 |
| | Do Vision and Language Models Share Concepts? A Vector Space Alignment Study <ul style="list-style-type: none"> • Measured the alignment of vector spaces in vision and language models with Procrustes. • Contributed with direction of data preparation, baselines, and paper writing. | TACL 2024 |
| | | |

| | |
|-------------------------|--|
| RESEARCH PROJECTS | Factual Consistency of Multilingual Language Models Findings of ACL 2022 |
| | <ul style="list-style-type: none"> Studied how consistent factual predictions are in different languages. Used a mixture of neural translation models to translate a resource of English paraphrases to 45 languages. |
| | Challenges and strategies in cross-cultural NLP ACL 2022 |
| | <ul style="list-style-type: none"> Surveyed and discussed strategies for cross-cultural and multicultural NLP. Contributed with the “Common ground” dimension that differs among cultures. |
| AWARDS AND HONORS | • Distinction , Graduated with the highest academic distinction 2019 |
| | • Dean’s list , All Semesters 2013-2018 |
| | • Scholar , IEEE Latin American Summer School in Computational Intelligence 2017 |
| | • 11th Place / 450 , Ascenci Dev Cup (France) Finalist 2015 |
| SKILLS | Languages: English, Spanish, French. |
| | Programming: Python, C++, SQL. |
| | Libraries: PyTorch, NumPy, HuggingFace, Pandas, Scikit-learn. |
| ACADEMIC SERVICES | Reviewers for: <i>ACL Rolling Review (ARR)</i> . |
| PUBLICATIONS | <ol style="list-style-type: none"> 1. Constanza Fierro, Negar Foroutan, Desmond Elliott, Anders Søgaard. How Do Multilingual Language Models Remember? Investigating Multilingual Factual Recall Mechanisms. Pre-print. 2. Constanza Fierro, Ruchira Dhar, Filippos Stamatiou, Nicolas Garneau, and Anders Søgaard. 2024. Defining Knowledge: Bridging Epistemology and Large Language Models. <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)</i>. Miami, Florida, USA. 3. Constanza Fierro, Reinald Kim Amplayo, Fantine Huot, Nicola De Cao, Joshua Maynez, Shashi Narayan, and Mirella Lapata. 2024. Learning to Plan and Generate Text with Citations. <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (ACL)</i>. Bangkok, Thailand. 4. Constanza Fierro, Nicolas Garneau, Emanuele Bugliarello, Yova Kementchedjhieva, and Anders Søgaard. 2024. MuLan: A Study of Fact Mutability in Language Models. <i>Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers) (NAACL)</i>, Mexico City, Mexico. 5. Jiaang Li, Yova Kementchedjhieva, Constanza Fierro, and Anders Søgaard. 2024. Do Vision and Language Models Share Concepts? A Vector Space Alignment Study. <i>Transactions of the Association for Computational Linguistics (TACL)</i>. 6. Constanza Fierro and Anders Søgaard. 2022. Factual Consistency of Multilingual Pre-trained Language Models. <i>Findings of the 60th Association for Computational Linguistics (ACL)</i>. Dublin, Ireland. 7. Daniel Hershcovich, Stella Frank, Heather Lent, Miryam de Lhoneux, Mostafa Abdou, Stephanie Brandl, Emanuele Bugliarello, Laura Cabello Piqueras, Ilias Chalkidis, Ruixiang Cui, Constanza Fierro, Katerina Margatina, Phillip Rust, and Anders Søgaard. 2022. Challenges and Strategies in Cross-Cultural NLP. In <i>Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (ACL)</i>. Dublin, Ireland. |